

COL 331 Minor 2

Viraj Agashe

TOTAL POINTS

36.5 / 50

QUESTION 1

1 Q1 10 / 10

Maintain a hierarchical structure for scheduling entity

✓ + 4 pts Use a 2-level CFS scheduler

✓ + 5 pts Correct explanation of the 2-level scheduler

+ 2 pts Incomplete explanation

+ 2 pts Partially correct explanation

✓ + 1 pts Saved the history per-user for the next scheduling

Other than the CFS algorithm for users

+ 2 pts Round-robin for selecting users

+ 2 pts Other than RR

+ 4 pts Correct explanation of the chosen algorithm

+ 2 pts partially correct

+ 2 pts Incomplete explanation

+ 1 pts Mentioned how processes are scheduled after selecting a user

+ 0 pts Incorrect or not attempted

QUESTION 2

2 Q2 0 / 10

✓ + 0 pts Not attempted/Wrong

+ 4 pts High-level idea (Using MLFQ, CFS)

+ 3 pts Details (Round-robin, datastruct)

+ 1 pts Time quanta

+ 2 pts Managing priorities

QUESTION 3

Q3 10 pts

3.1 (a) 3 / 3

✓ + 1 pts both reference the same list

✓ + 1 pts sigpending structure correctly described

✓ + 1 pts sigqueue structure correctly described

+ 0 pts incorrect/unattempted/unclear

3.2 (b) 1.5 / 3

✓ + 1.5 pts return address set to the instruction immediately following the instruction that was interrupted by the signal

+ 1.5 pts return address including other cpu state pushed on stack

+ 0 pts incorrect/unattempted

3.3 (c) 3 / 3

✓ + 0.5 pts resource grant clause

✓ + 0.5 pts inheritance clause

✓ + 2 pts sound arguments

+ 1 pts unclear/incomplete arguments

+ 0 pts incorrect/unattempted

3.4 (d) 1 / 1

✓ + 1 pts correct answer

+ 0 pts incorrect/unattempted

QUESTION 4

4 Q4 10 / 10

✓ - 0 pts *Correct*

- 10 pts Incorrect

- 2 pts Semantic errors, (like use or instead of
and or incorrectly addressing an array etc)

- 4 pts Partial solution/Inefficient/Leads to
deadlock/Pseudocode not provided

- 6 pts Right track but missing a more detailed
explanation/Pseudocode

- 3 pts Needs more detailed explanation (How
are the semaphore used, which forks to choose
etc)

QUESTION 5

Q5 10 pts

5.1 (a) 5 / 5

✓ + 5 pts *Correct*

+ 3.5 pts Partially correct/ Diagram missing

+ 0 pts Incorrect/ Not attempted

5.2 (b) 3 / 5

+ 0 pts Incorrect/ Not attempted

+ 5 pts Correct

✓ + 3 pts *Partially Correct*

COL331/COL633
OPERATING SYSTEMS
MINOR-2

TIME: 1 HOUR

MARKS:

50

INSTRUCTIONS:

1. All the questions are **compulsory**.
2. No doubts will be addressed during the exam. Make your *assumptions*.

Questions:

1. How can we modify the CFS scheduling policy to *fairly* allocate processing time among all users instead of processes? Assume that you have a single CPU and all the users have the same priority (they have an equal right to the CPU regardless of the processes that they spawn). Each user may spawn multiple processes, where each process will have its individual CFS priority between 100 and 139. Do not consider the real-time or deadline scheduling policies.

[10 Marks]

Let us say the users are u_1, u_2, \dots, u_n . Instead of ensuring fair time amongst processes, we need to ensure among users. For this, we can allocate CPU time (scheduling slice) on the basis of the cumulative v-runtimes of the processes within a user's processes.

~~Let us say we have a strict user & ? for each user of the system.~~ We will maintain a red-black tree containing the v-runtimes (cumulative) of all processes owned by ~~all the users~~ by each user.

At any time, we select the user with the minimum cumulative v-runtime of all its processes and allot it a time slice. ~~We can calculate the scheduling slice on the basis of the cumulative v-runtime similar to the common fair scheduler.~~ Similarly to CFS, we allot a fixed time slice to each user since priority of each user is the same (~~the effect of priorities~~ (scheduling - slice)).

Name: _____

Entry Number: _____

Now within a user's processes we can simply use the CFS algorithm as before — pick an individual task with the currently min vruntime & schedule it.

Use the priorities of the process to appropriately scale its vruntime to get the v-runtime, and once it is pre-empted, add this v-runtime to the cumulative v-runtime of the user.

In this way we can have a CFS algorithm for users.

2. Given a mixture of interactive, I/O intensive and long running processes whose execution time is not known a priori, design a scheduling algorithm for a single CPU that optimizes the completion time as well as the responsiveness of the interactive jobs. The algorithm (and associated data structures) should take into account the diversity of jobs and the fact that new jobs and high priority jobs need quick service, whereas low-priority long-running batch jobs can be delayed (read deprioritized).

[10 Marks]

Name: NIRAJ AGASME

Entry Number: _____

3. Answer the following questions about signal handling and real-time scheduling.
[3 + 3 + 3 + 1 Marks]
- a. Do *struct sigpending* and *struct sigqueue* reference the same data structure? Explain.

Name: _____

Entry Number: _____

~~struct~~ sigpending is essentially a linked list of sigqueue objects.



The list head of sigpending points to the next sigqueue item in the linked list, while the list head of each sigqueue object points to the next & prev. elements in the sigpending linked list.

b. How is the return address of a signal handler set? What is it set to?

The return address of a signal handler is set to the address of a glibc function (essentially a system call) which restores back the state (context) of original process. It is set by calling the sigreturn instruction.

c. Prove that in PCP, once the first resource is acquired, there can be no more priority inversions (provide a very short proof).

In PCP we know that a process ~~can acquire a resource if~~ has ownership of a resource if:

- It was the one that set the system ceiling
- It has a prio genuinely higher than the system ceiling.

Now, consider the resource R, under contention. Say that a process P_1 is holding it, and P_2, \dots, P_k are waiting. We claim that $\text{prio}(P_1) \geq \text{prio}(P_j) \forall j$.

If this is not true, then \exists some P_j s.t. $\text{prio}(P_j) > \text{prio}(P_1)$.

However in the PCP algorithm we set $\text{prio}(P_1)$ to the max. of the waiting tasks (soft-increase). Further, when R was acquired, by the

conditions above, P must have set the system ceiling, so it was genuinely the highest prio process. \Rightarrow No priority inversions.

d. For a system with periodic and pre-emptive jobs, what is the utilization bound (maximum value of U till which the system remains schedulable) for EDF? Ans:

1

4. Write pseudocode (C like syntax) for solving the Dining Philosopher's Problem. Use only semaphores. Use three states for each philosopher: THINKING, HUNGRY, and EATING. [10 Marks]

Sol We can solve the dining philosopher's problem as follows:

semaphore S_1, S_2, \dots, S_5 ;

~~uses~~ philosophers P_1, P_2, \dots, P_5 ;

~~if~~ // Let rank denote which philosopher I am.

~~If (rank == P5) {~~
~~wait(S4);~~
~~wait(S5);~~
~~state == EATING;~~
~~...~~
~~state == THINKING;~~
~~post(S4); post(S5);~~
~~}~~

If (rank == P5 && state == HUNGRY) {


wait(S4);
 wait(S5);
 state == EATING;

...
 state == THINKING;
 post(S4); post(S5);
 }


else { if state == HUNGRY } {

wait(Si+1);
 wait(Si-1);
 state == EATING;

...
 state == THINKING;
 post(Si+1); post(Si-1);
 }



One of the philosophers try left
 Others all try right
 fork



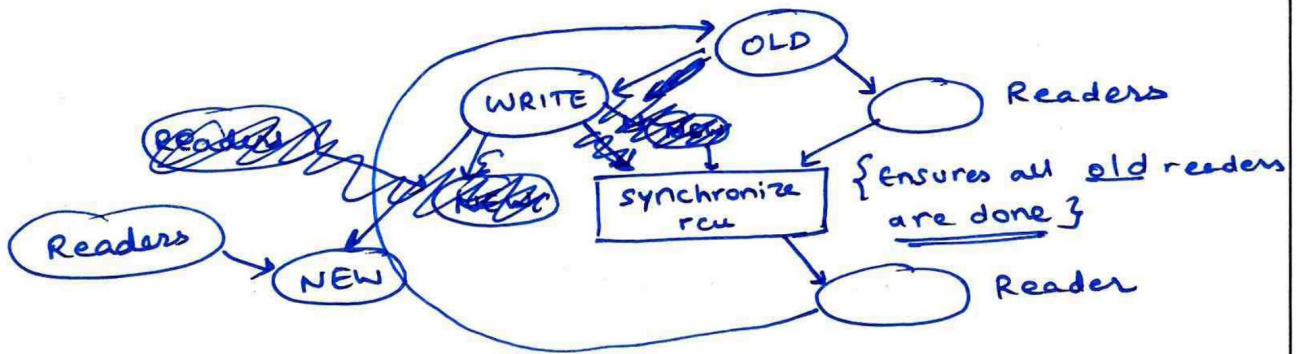
Name: _____

Entry Number: _____

5. Answer the following questions regarding RCU:

- a. Prove that no reader can be alive (reading the old value of the pointer) when `synchronize_rcu` returns. Create diagrams with happens-before edges and prove that such a situation is not possible. Show a proof by contradiction. [5+5] marks.

We know that `synchronize_rcu` ensures that all readers have exited their RCU protected regions \Rightarrow Nobody can be reading that pointer at this time. So we use the following graph—



Consider the above graph of happens-before relationships. Suppose that earlier readers dereferenced the old value of a pointer, which was written sometime. Any reader called after `synchronize_rcu` has also been called after the old value has been deleted. Reading the old value \Rightarrow we read it before the write. This however creates a cycle in our happens-before graph, which is not possible \Rightarrow No reader can see old value after `rcu-synchronize` has been called.

- b. Show the pseudocode for registering and deregistering readers and `synchronize_rcu` with a real time preemptive kernel.

Register.

```
register () {
    CPU-Reading = 1
    CPU-Register-Reader ();
}
```

This function will set the bit that the CPU is in a RCU protected region in the VEB tree.

deregister {

```
    CPU-deregister-reader ();
}
```

This is a callback function which sets the bit in the VEB to 0 again.

Synchronize-rcu {

bool b = check_bits(); —————→ Check if all bits
in VEB are 1
(check root)

if (!b) {

~~bitmask~~ ~~bm~~ ~~=~~ ~~get-VEB-bits~~);

~~while (bm) { count = 0;~~

~~if (bm & 0xffff0) {~~

~~send-ipi(count); // Send IPI to "count"~~
~~CPU (index into the~~
~~VEB tree)~~

~~average bit error~~

count = 0;
while (count < ncpus) {
if (VEB-get(count) == 0) {

send-IPI(count); // Send an IPI to
} else { count++; continue; } "count"-th processor(cpu)
~~while~~ ~~wait~~ ~~on~~ (VEB-get(count) == 1); calls the defer function
(callback).

~~for~~ while (VEB-get(count) == 1);

Wait till
bit is set to
1